

Final_Project_supervised_learning

March 7, 2025

```
[162]: #run the next cell in case certain libraries are missing
```

```
[163]: #!pip install nltk  
#!pip install powerlaw
```

```
[164]: import numpy as np  
import matplotlib.pyplot as plt  
import powerlaw  
import scipy.stats as stats  
from sklearn.linear_model import LogisticRegression  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score, classification_report  
import nltk  
from collections import Counter  
from nltk.corpus import brown  
from nltk.corpus import gutenberg
```

```
[304]: #training corpus to gauge fit / p-value  
#gutenverg corpus that is primarily full books, stories, and American classics  
#medium p value > 0.5, however, it varies considerably from the p value of the  
→brown corpus, which is almost 1  
#implying potentially different writing styles  
  
def get_corpus_0_training():  
    nltk.download('gutenberg')  
    lst = gutenberg.words()  
    return lst, "gutenberg"  
  
lst, lb = get_corpus_0_training()
```

```
[nltk_data] Downloading package gutenberg to /home/jovyan/nltk_data...
```

```
[nltk_data] Package gutenberg is already up-to-date!
```

```
[305]: #given a list of strings, identfiy and remove those that are not words.  
def tokenize(lst):
```



```

plt.ylabel("Word Frequency")
#lets do a log graph
plt.xscale("log")
plt.yscale("log")
diff_words = len(frequencies)
#avg_frequency = total_words/diff_words
#print("number of different words", diff_words)
print("\n ")
return plt

# a function to show the sequence of ratios between the most frequent words
def ratio_plot(rf_tuple):
    ratio_lst = []
    for i in range(0,len(rf_tuple)-1):
        ratio = rf_tuple[i][1] / rf_tuple[i+1][1]
        ratio_lst.append(ratio)
    #print(ratio_lst)
    frequencies = []
    plt.figure(figsize=(10,4))
    plt.plot(ratio_lst)
    plt.xlabel("sequence iteration")
    plt.ylabel("ratio of sequence pair")

#uses maximum likelihood estimation to determine ideal s_mle coefficient
def KS_test(rf_tuple):
    #print(lb)
    ranks,freqs = zip(*rf_tuple)
    ranks = np.arange(1, len(rf_tuple) + 1)
    i=0
    s_evol=[]
    print("Using Maximum Likelihood Estimation to find best coefficient")
    for i in range(2,len(freqs)):
        m_freqs = np.array(freqs[:i])
        fit = powerlaw.Fit(m_freqs, verbose=False)
        s_mle = fit.alpha
        s_evol.append(s_mle)
    #print(s_evol)

    print(f"Estimated Zipf exponent using MLE: {s_mle:.4f}")

    #hyper_parameters
    s = 1
    #print(freqs)
    C = freqs[0]
    #print(C)
    predicted_freqs = C / ranks**s

```

```

observed_cdf = np.cumsum(freqs) / np.sum(freqs)
predicted_cdf = np.cumsum(predicted_freqs) / np.sum(predicted_freqs)

ks_val, p_val = stats.ks_2samp(observed_cdf,predicted_cdf)

print(f"KS Statistic: {ks_val:.4f}")
print(f"P-value: {p_val:.5f}")

def gen_data():
    #print(rf_tuple)
    adj_val = 100
    #this can be scaled depending on the size of the total dataset
    #100 happens to be a fairly consistent value regardless of the dataset size
    plot(rf_tuple[:adj_val])
    ratio_plot(rf_tuple[:adj_val])
    print(lb)
    KS_test(rf_tuple[:adj_val])

```

```
[307]: gen_data()
```

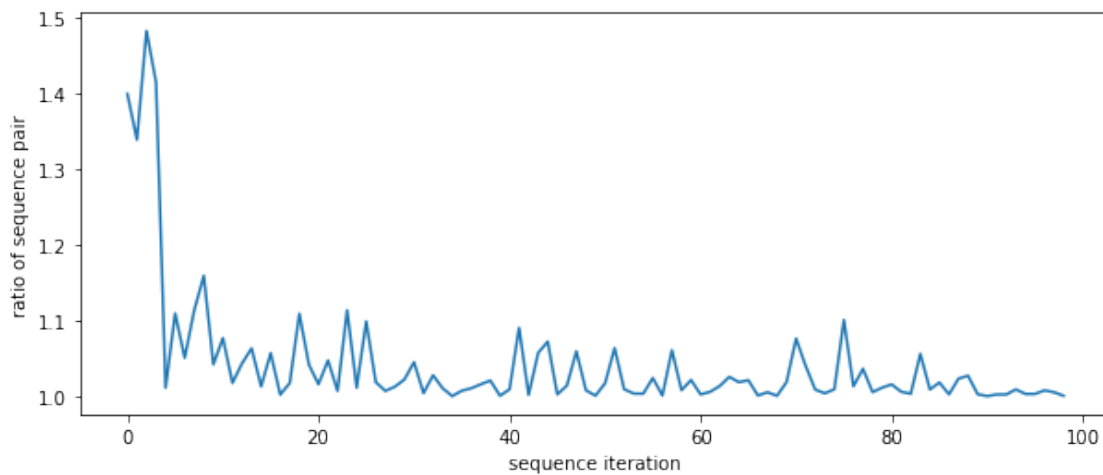
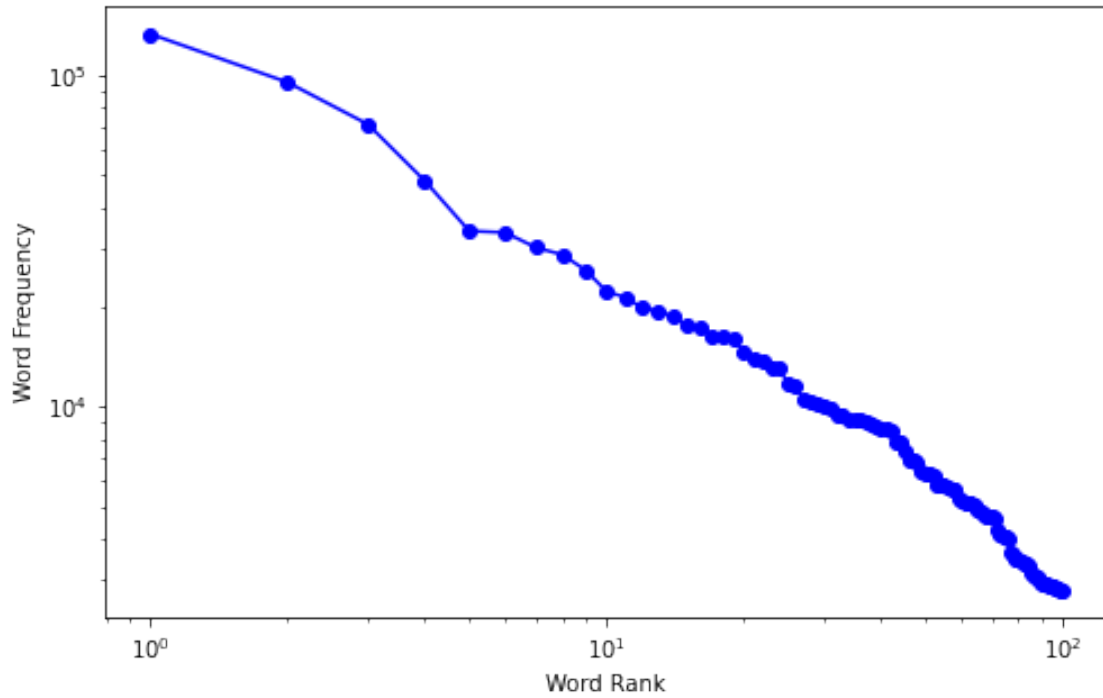
gutenberg

Using Maximum Likelihood Estimation to find best coefficient

xmin progress: 65%Estimated Zipf exponent using MLE: 2.4056

KS Statistic: 0.1100

P-value: 0.58301



```
[300]: #training corpus to gauge fit / p-value
#brown corpus has a wide variety of writing that closely resembles everyday
↳ speech and writing

def get_corpus_1_training():
    nltk.download('brown')
    lst = brown.words()
```

```
return lst,"brown"
```

```
lst,lb = get_corpus_1_training()
```

```
[nltk_data] Downloading package brown to /home/jovyan/nltk_data...
```

```
[nltk_data] Package brown is already up-to-date!
```

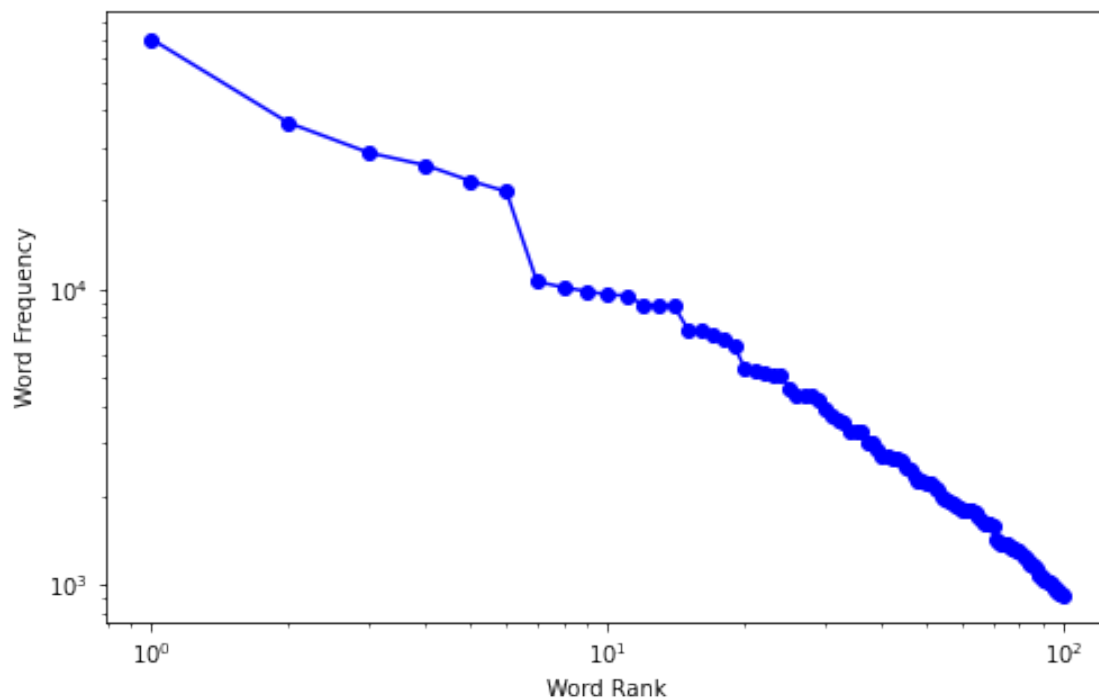
```
[261]: gen_data()
```

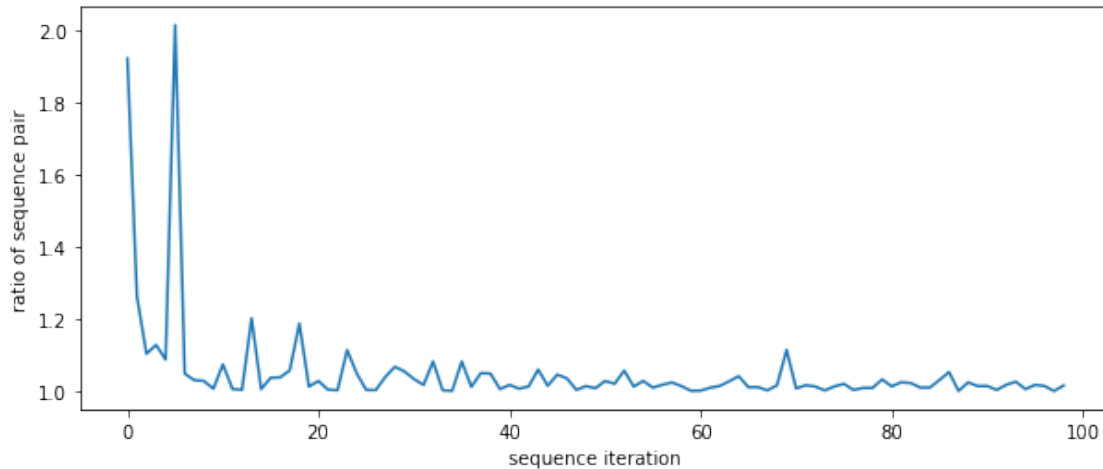
brown

Using Maximum Likelihood Estimation to find best coefficient

Estimated Zipf exponent using MLE: 1.9942

P-value: 0.99969





[262]: #observe the two graphs

#the first one is the brown corpus, which has about 1 million words
 #the second is the gutenber corpus which has about 2 million words

#for both datasets, we only take the first 20->100 words or so, as the tail of
 ↳the rank frequency distribution can skew
 #the zipfian distribution

#decreasing the size of the frequency list can be seen as decreasing the
 ↳variance while increasing bias

#notice how even the brown corpus is less than size the gutenber corpus, it is
 ↳better fitted using the KV test,
 #and shows a stronger correlation to zipf's law.
 #This is very interesting because typically increasing the size of the text
 ↳would show a stronger zipfian distribution

[263]: #the ratio of sequence pairs

#is an experiment to see how much each word differs from the next.
 #we do not use this graph to train our model, but it remains an interesting
 ↳sub-topic to delve into
 #notice the distribution of the frequencies for both the gutenber and the
 ↳brown corpi have a very similar pattern

[69]: #we deduce that merely the number of words is not the only factor, but the
 ↳differing writing styles
 #which leads us to design a hypothesis, do different writing styles show
 ↳different p values when fitted using the power law?

```
#gutenberg corpus: focus on literature, classics, full books and stories
#brown corpus: larger variety of text from a variety of genres, including news
↳ articles, nonfiction and fiction

#We suspect that high p values correspond more to everyday text while lower p
↳ values correspond more to artistic writing
#We now find our test sets.
```

```
[287]: gen_data()
```

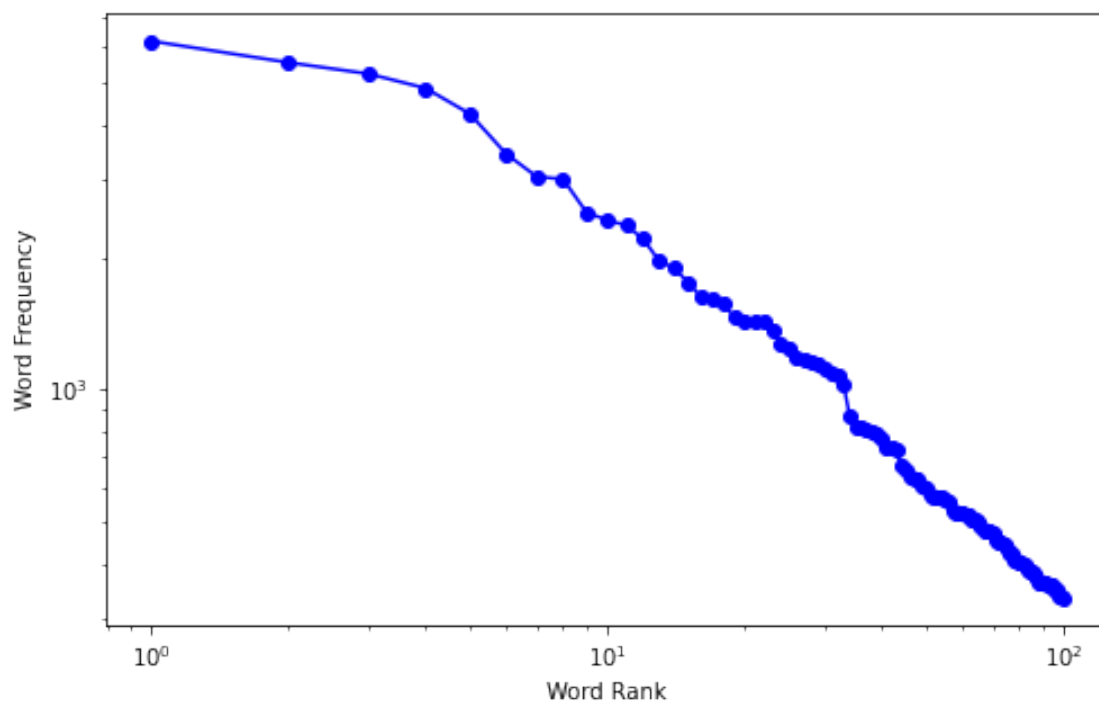
shakespeare

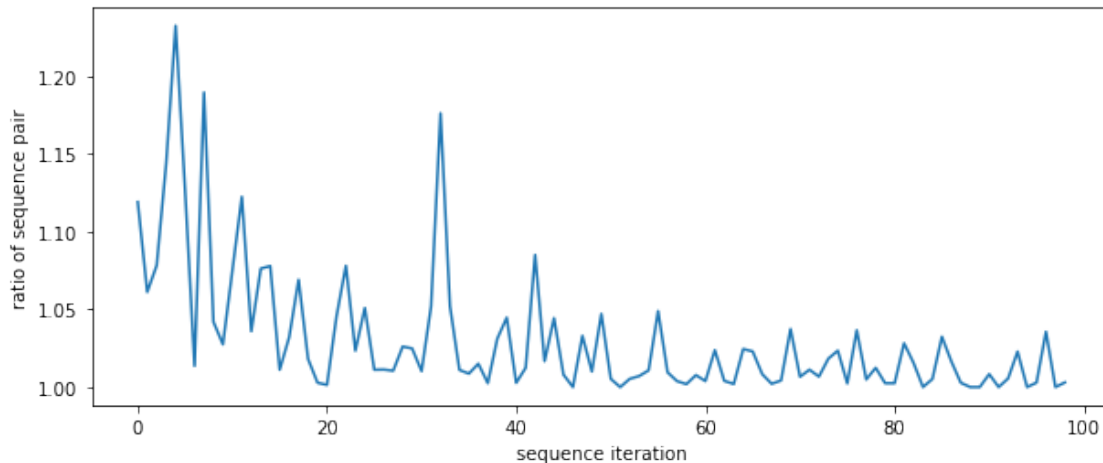
Using Maximum Likelihood Estimation to find best coefficient

Estimated Zipf exponent using MLE: 2.1927

KS Statistic: 0.1600

P-value: 0.15484





```
[284]: #test corpus of shakespeare's plays
#suspected to have a low p_value given shakespeare's very wide vocabulary

from nltk.corpus import shakespeare

def get_corpus_2_testing():
    nltk.download('shakespeare')
    lst = []
    for file_id in shakespeare.fileids():
        lst.extend(shakespeare.words(file_id))
    #print(lst)
    return lst, "shakespeare"

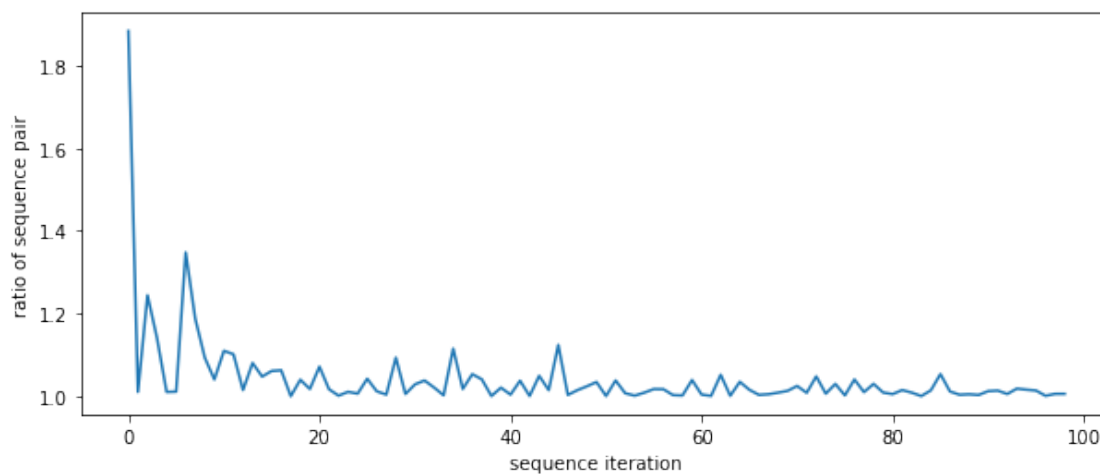
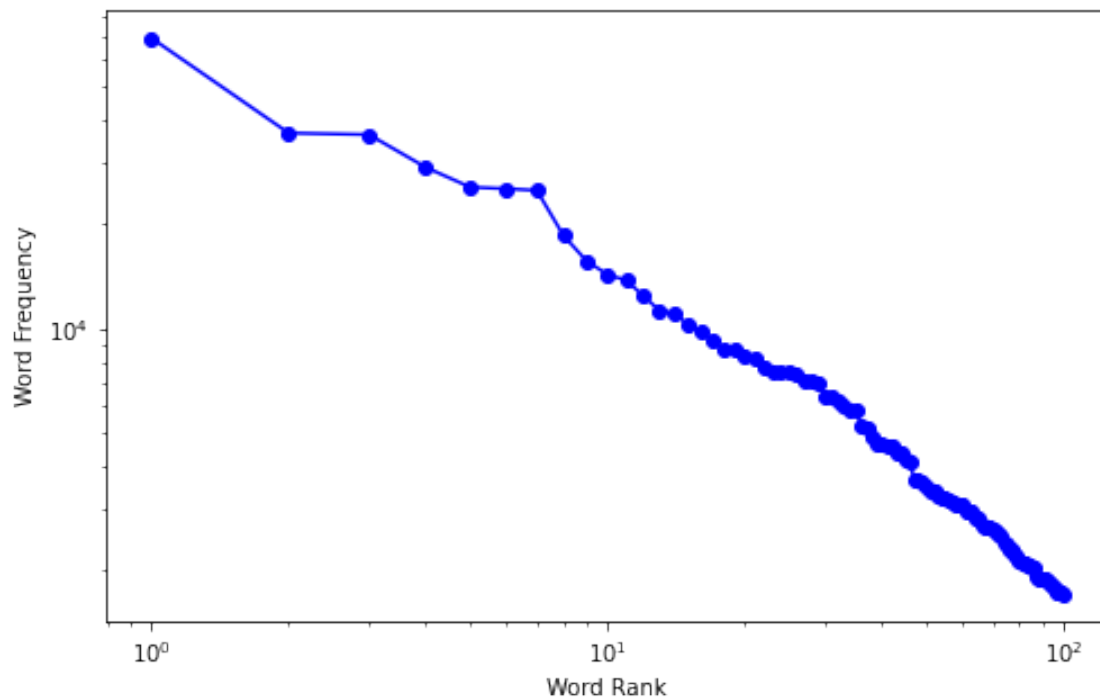
lst, lb = get_corpus_2_testing()
```

```
[nltk_data] Downloading package shakespeare to
[nltk_data] /home/jovyan/nltk_data...
[nltk_data] Package shakespeare is already up-to-date!
```

```
[291]: gen_data()

#as suspected, Shakespeare follows zipfs law less closely with a low p_value
```

```
reuters
Using Maximum Likelihood Estimation to find best coefficient
Estimated Zipf exponent using MLE: 2.0327
KS Statistic: 0.1100
P-value: 0.58301
```



```
[295]: #test corpus of news articles
#suspected to have a high p_value
#while the p value was not as high as we suspected, it was still significant_
→ enough to determine a zipfian distribution

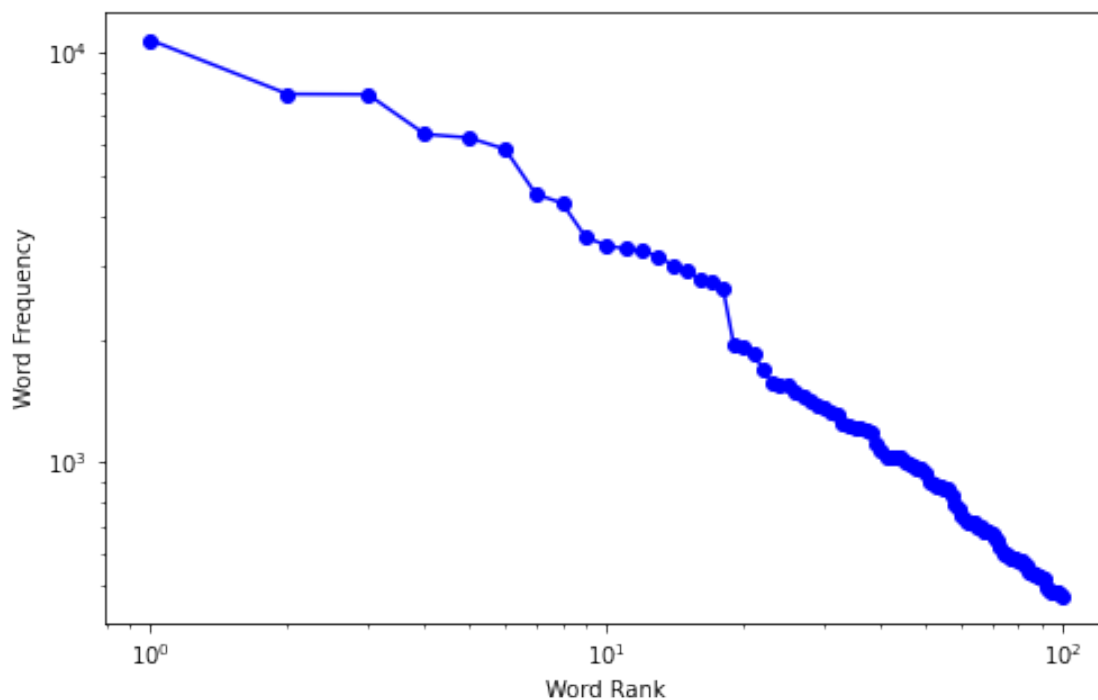
from nltk.corpus import reuters
```

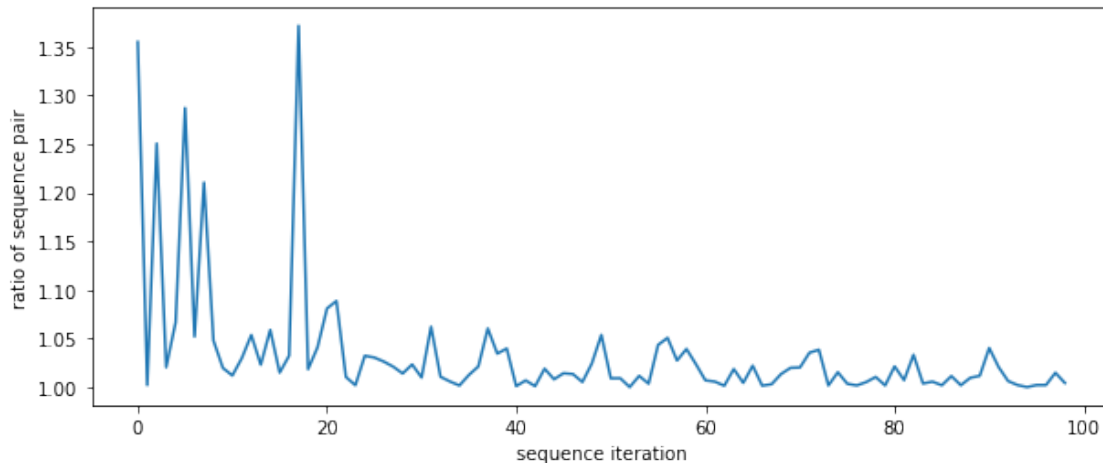
```
def get_corpus_3_testing():  
    nltk.download('reuters')  
    lst = reuters.words()  
    return lst, "reuters"  
  
lst, lb = get_corpus_3_testing()
```

```
[nltk_data] Downloading package reuters to /home/jovyan/nltk_data...  
[nltk_data]   Package reuters is already up-to-date!
```

```
[299]: gen_data()
```

webtext-poetry
Using Maximum Likelihood Estimation to find best coefficient
Estimated Zipf exponent using MLE: 2.1829
KS Statistic: 0.1600
P-value: 0.15484





```
[104]: # based on our assumptions, we see that a p value > .4 implies a more
      ↪ consistent, easy to gauge pattern in writing that
      # we associate with more everyday speech

      #a low p_value < .2 implies a less consistent ptern that we associate with
      ↪ more elaborate writing, like in poetry or plays
```

```
[296]: #test corpus of webtext poetry
      #suspected to have a low p (< 0.2)

      from nltk.corpus import webtext

      def get_corpus_4_testing():
          nltk.download('webtext')
          lst = webtext.words()
          return lst, "webtext-poetry"

      lst,lb = get_corpus_4_testing()
```

[nltk_data] Downloading package webtext to /home/jovyan/nltk_data...

[nltk_data] Package webtext is already up-to-date!

```
[147]: #test_corpus_2 (shakespeare) was suspected to not follow zipfs law because
      #which is confirmed with the p value it generates

      #test_corpus_3 (retuers) was suspected to follow zipfs law
      #this is confirmed with the relatively high p value it generates

      #test_corpus_4 (webtext(poetry)) was suspected to not follow zipfs law
      #this is confirmed with the relatively small p value it generates
```

[148]: *#this classification technique perhaps allows us to create a spectrum of*
→writing styles. Perhaps genres that are more ambiguous follow zipf's law more
#for this specific project, we attempted to find 2 extremes
#given the subjectivity involved in determining genre, there is room for error
→depending on the text compiled
#with that said, we can see a significant deviation in p value that we can at
→least partly attribute to writing style

#another thing to keep in mind is the generated s_coeff value - lower
→s_coefficient values are correlated with higher p values

[229]: *#conclusion and further discussions*

#Using a combination of techniques (logistic regression, maximum likelihood
→estimation)
#and several principles (zipf's law, ratio sequence)
#we can use the p value of an attempted fit in order to infer the type of
→writing styles involved in a text
#given that writing style, like a p value, can be understood as a spectrum, we
→attempted to find writing styles and combinations
#that gave high and low p-values, respectively

#for efficiency and estimation reasons, we decided to look only at the
→distribution of the first 100 words

#we learned that poetry and plays tend to have a low p value of 0.1, while the
→reuters article had a p value of approximately 0.5
#and the brown corpus had a very high p value of .9998

#the gutenbergs corpus we used to train our concept model had a compilation of
→different texts, but they varied less than that of the brown corpus
#increasing the variation of text is more likely to cause a stronger zipfian
→distribution.

#other thoughts:
#graphing the sequence of ratios for nearly every data set showed a striking
→pattern, where as the number of different words we
#looked at increased, the sequence of words that had an identical frequency
→(ratio of 1) became longer.

#elaborations and further experimentation
#other ideas that I played around with was measuring the stability of the s
→value as the number of different words increased
#and different means of collecting data

```
#one idea is to use corpi where each corpus was written by a single author,␣  
→instead of attempting to determine the specific genre  
#the problem with this approach is finding enough authors who write␣  
→prolifically so our dataset is large enough
```

```
#notice how we did not the words themselves as a criteria for classification
```

[230]: #references

```
#this project was my own, but the concept of zipfs law and the␣  
→Kolmogorov-Smirnov test was studied  
#in order to approach this problem  
  
#all code is my own but many python libraries were used
```